

Determining Ramanujan numbers

Arjen Markus

April 11, 2024

Introduction

As described at the MathWorld webpage <https://mathworld.wolfram.com/TaxicabNumber.html>, Ramanujan or taxicab numbers come in two flavours, both related to the sum of two cubic integer numbers. In this note I follow the definition given in Sloan's on-line encyclopedia of integer sequences, <https://oeis.org/A001235>: a Ramanujan number is an integer that can be expressed as the sum of two cubic numbers in at least two ways. The smallest example is 1729:

$$1729 = 1^3 + 12^3 = 10^3 + 9^3 \quad (1)$$

The first few are (copied from the OEIS site): 1729, 4104, 13832, 20683, 32832, 39312, 40033, 46683, 64232, 65728, 110656, 110808, 134379, 149389, ...

It is easy enough to determine them with brute force and I do not know if any other methods are known. But that is not the point here. This note explores a few implementations to determine them. The idea is simple: let i and j run over the interval $[1, n]$ and calculate the sums of $i^3 + j^3$ for some suitable value of n . If the values occur more than once, we have found a Ramanujan number. Also note that interchanging i and j does not give a new value, so i should run over the interval $[1, j]$.

Side note: I first let the integers run over the interval $[0, n]$, but then I realised that a Ramanujan number with one of the two being zero would amount to a non-trivial solution to Fermat's last theorem. So I dropped that case.

Plain and simple

Here is a program that takes a straightforward approach:

```
program ramanujan_plain
  implicit none

  integer, parameter :: max_ij = 100
  integer, parameter :: max_sum = 2 * max_ij ** 3
  integer              :: cnt(max_sum)
```

```

integer                :: i, j, k

!
! Determine the sums of two cubes
!
cnt = 0

do j = 1,max_ij
  do i = 1,j
    k = i ** 3 + j ** 3

    cnt(k) = cnt(k) + 1
  enddo
enddo

!
! Write out the numbers for which there are at least two
! combinations i,j that give these numbers
!
do k = 1,max_sum
  if ( cnt(k) > 1 ) then
    write(*,*) k
  endif
enddo
end program ramanujan_plain

```

I chose n to be $100 - \text{max_ij}$ in the program, so that the maximum value of the sum is 2 million.

The program simply loops over the two integers and counts how often the value occurs. For this, it uses a large array, `cnt`, plain and simple. Then it only prints those values that occur more than once.

The outcome is:

```

1729
4104
13832
20683
32832
39312
40033
46683
64232
65728
110656
110808
134379

```

149389
165464
171288
195841
216027
216125
262656
314496
320264
327763
373464
402597
439101
443889
513000
513856
515375
525824
558441
593047
684019
704977
805688
842751
885248
886464
920673
955016
984067
994688
1009736
1016496

You would expect to see several numbers beyond 1 million, given that they are not too sparse, the intervals between successive numbers are rather small. The largest one found is much smaller than the limit of 2 million. This has to do with the fact that a value below 2 million may be the result of one of the integers being larger than 100 and the other smaller. So we do not find all Ramanujan numbers smaller than 2 million. You could adapt the program to take care of this, but it would make it less regular. And that is what I want to explore.

Using array operations

If we use Fortran's array operations, then we can make the program quite a bit shorter (leaving out the declarations for brevity, they are more or less the same

as in the first version):

```
!
! Determine the sums of two cubes
!
sums = [(i**3+j**3, i = 1,j), j = 1,max_ij)]
cnt = 0
cnt(sums) = cnt(sums) + 1

!
! Write out the numbers for which there are at least two
! combinations i,j that give these numbers
!
write(*,*) pack( [(k, k = 1,max_sum)], cnt > 1)
```

The first line constructs an array via two implied-do loops and using the automatic reallocation feature, so we do not have to allocate the array `sums` beforehand and worry about the precise size.

Then to determine how often a particular sum occurs, we use a vector index:

```
cnt(sums) = cnt(sums) + 1
```

The accumulation occurs without us having to program that explicitly.

Finally, to determine which values are indeed Ramanujan numbers we use the `pack()` function. Because it only returns the elements of the array that is being packed, not the indices of the elements that are retained, we need to construct an array of such indices, but that is easily done.

(The outcome of this program is the same as the first one, except that several numbers are printed on the same line. That detail can be fixed of course.)

A compact version

A drawback of both these versions is that they rely on a large array to store the counts. If we want to determine larger and larger Ramanujan numbers, then this becomes a burden. So, here is an alternative: determine only the sums and count per value how often it occurs, then store the Ramanujan numbers in a separate array.

The code can be fairly compact as well, though we need an explicit do-loop, but the memory use is greatly reduced:

```
program ramanujan_compact
  implicit none

  integer, parameter :: max_ij = 100
  integer, allocatable :: saved_sums(:)
  integer, allocatable :: sums(:)
```

```

integer                :: i, j, value

!
! Determine the sums of two cubes
!
sums = [((i**3+j**3, i = 0,j), j = 0,max_ij)]
allocate( saved_sums(0) )

do while ( size(sums) > 0 )
    value = sums(1)

    if ( count( sums == value ) > 1 ) then
        saved_sums = [saved_sums, value]
    endif

    sums = pack( sums, sums /= value )
enddo

!
! Write out the numbers for which there are at least two
! combinations i,j that give these numbers
!
write(*,*) saved_sums
end program ramanujan_compact

```

The program first determines the sums and then examines the values one by one. The Ramanujan numbers are stored and in any case the elements of the array `sums` that are equal to the first value are removed for the next round. Thus a rough estimate of the memory use is at most three times the number of sums ($\approx 3n^2/2$) – once for the sums, once for the logical mask and once for the Ramanujan numbers, instead of $2n^3$. Numerically: 15,000 versus 2 million.

A surprise, perhaps, is that the numbers are not printed in a sorted order (edited to fit on a page):

1729	4104	13832	20683	32832	46683
39312	40033	65728	64232	110656	134379
110808	165464	149389	216125	171288	195841
327763	216027	262656	373464	402597	314496
320264	525824	439101	593047	515375	513856
558441	443889	513000	684019	885248	842751
704977	955016	886464	1016496	805688	984067
994688	1009736	920673			

Reduce it to a simple statement

The implementations so far consist of two steps: determine the Ramanujan numbers and then print them. With a bit of ingenuity (the program is not easy to read) you can combine these two steps and end up with a program that essentially consists of one composite statement (leaving out the trivial declarations):

```
write(*,*) pack( [ (k, k = 1,max_sum)], &
                  [ ( count( [(i**3+j**3, i = 1,j), j = 1,max_ij]) &
                        == k ) > 1, k = 1,max_sum )] )
```

Because it computes the set of sums over and over again (two million times) it is very slow. If you store the results of the double implied do-loop, then it is quite a bit faster.

The program seems a challenge for some compilers. This is because compilers will try and evaluate the statement at *compile time* and that seems rather compute-intensive. In fact, the maximum n (or `max_ij` in the program) directly influences the time it takes for the compilation to finish, which would be very unlikely if the statement was treated as any ordinary code that should be run at run-time.

Conclusion

The excuse for this note was a mathematical topic that simply intrigues me, but the real subject is the variety of implementations that are possible, each having its own pros and cons, in terms of length, readability and memory use. Important instruments in the code presented are the array operations and the `pack()` function.